

Introduction to Algorithms and Data Structures

Lecture 8: Sets, dictionaries and hashing

John Longley

School of Informatics
University of Edinburgh

October 2021

Sets and dictionaries

Two important datatypes ...

- ▶ (Finite) sets of items of a given type X . E.g. $\{3, 5\} = \{5, 3\}$

contains : $X \rightarrow \text{bool}$
insert : $X \rightarrow \text{void}$
delete : $X \rightarrow \text{void}$
isEmpty : $\text{void} \rightarrow \text{bool}$

- ▶ Dictionaries (i.e. lookup tables) mapping keys of type X to values of type Y .

lookup : $X \rightarrow Y$
insert : $X * Y \rightarrow \text{void}$
delete : $X \rightarrow \text{void}$
isEmpty : $\text{void} \rightarrow \text{bool}$

Sets and dictionaries in Python



```
Beatles = {'John', 'Paul', 'George', 'Ringo'}  
'George' in Beatles           # returns True
```

```
BeatlesYearsOfBirth =  
    {'John':1940, 'Paul':1942, 'George':1943, 'Ringo':1940}  
BeatlesYearsOfBirth['George']    # returns 1943
```

Sets and dictionaries via sorted arrays

Could implement sets/dictionaries via (any impl of) lists:

```
Beatles_Rep = ['John', 'Paul', 'George', 'Ringo']
```

```
BeatlesYearsOfBirth_Rep = [('John', 1940), ('Paul', 1942), ...]
```

But average-case time for **contains/lookup** will be $\Theta(n)$ (**terrible!**)

Much better if arrays are **sorted (by key)**.

Can then use **binary search**. E.g. for dictionaries:

```
binarySearch(A, key, i, j):    # searches A[i], ..., A[j-1]
    if j-1 = i
        if A[i].key = key then return A[i].value else FAIL
    else
        k =  $\lfloor i+j/2 \rfloor$ 
        if key < A[k].key then return binarySearch(A, key, i, k)
        else return binarySearch(A, key, k, j)
```

Using this, **contains/lookup** have worst-case time $\Theta(\lg n)$.

But **insert/delete** still costly. **Can we do better?**

Hash tables

Suppose our keys are strings (e.g. people's names). Number K of **potential** keys is vast — number n of **actual** keys 'currently in use' much smaller.

Really silly idea: Give a way of converting strings s to integers $\iota(s)$ (E.g. treat ASCII characters as digits to base 128). Then store value associated with s in a big **array** at position $\iota(s)$.

Impractical: K normally far too large, and most of the array would be unused.

More sensible idea: Choose some **hash function** $\#$ mapping potential keys s to integers $0, \dots, m-1$ (**hash codes**), where $m \sim n$.
Want $\#$ to be **easy to compute**. E.g. we might define:

$$\#(s) = \iota(s) \bmod m$$

Then **try to** use an array A of size m , storing the entry for key s at position $\#(s)$ in A .

Hashes and clashes

Problem: What if $\#(s) = \#(t)$ for two keys s, t ?

How likely are clashes to arise? E.g. if we took e.g. $m \sim 5n$ (and accepted the space wastage), would clashes be improbable?

Example: Keys are people, $m = 366$, $\#(p) = \text{birthday of } p$.

How many people must there be for probability of shared birthday to be $> 1/2$? (Assume uniform distrib.)

Answer: Just **23**! (Sometimes called the **birthday paradox**.)

See CLRS 5.4.1 for analysis (if you're interested).

Question: In a class of 347 (assuming uniform distrib), what would be the probability of a birthday shared by 2 people? By 3 people? By 4, 5, 6, 7, ...?

2	3	4	5	6	7
$> (100 - 10^{-123})\%$	$> 99.9999\%$	$> 99.8\%$	66%	15%	2%

The Great Inf2 Birthday Experiment (**totally optional!**)



Curious about shared birthdays?

Email j.r.longley@ed.ac.uk to take part!

Subject line: birthday **MMDD**

(E.g. birthday 0216 for 16 Feb; birthday 1006 for 6 Oct.)

Please observe this format: lowercase birthday then single space.

Body text: Your name if you'd like to be put in touch with any others with same birthday. Otherwise blank.

Closing date: 12 noon (BST) on **Friday 22 October**.

All emails/records will be deleted by Friday 29 October.


I'll announce the results on Piazza: number of participants, dates of any shared birthdays, and number of people sharing them.

I won't announce any names.

Dealing with clashes

So we must accept clashes (a.k.a. **collisions**) as a fact of life.

Solution 1: Store a list of entries (or **bucket**) for each hash value.



(Omit **value** components if it's just a set.)

Write n for number of entries, m for array size.

The ratio $\alpha = n/m$ is called the **load** on the hash table:
may be ≤ 1 or > 1 .

If we've decided on a desired load α , can 'expand-and-rehash' any time n gets too large (amortized cost is reasonable).

Bucket-list hash tables: some analysis

Recall: n table entries, m hash codes, $\alpha = n/m$.

Write b_i for number of entries in i th bucket.

Let's analyse average time for an **unsuccessful** lookup.

Assume that for k not in the table, $\#(k)$ equally likely to be any of the m hash codes.

If $\#(k) = i$, lookup will do b_i key comparisons if unsuccessful.

So **average** number of key comparisons is

$$\frac{1}{m} \sum_{i=0}^{m-1} b_i = n/m = \alpha$$

If computing $\#(k)$ itself takes $O(1)$ time, conclude that average time for unsuccessful lookup is $\Theta(\alpha)$. (Thinking of $\alpha \rightarrow \infty$.)

Can also show the same for **successful** lookup, assuming all keys present in table are equally likely. See CLRS 11.2.

Making a *proper* hash of it

Rarely true that all *potential* keys (e.g. strings) ‘equally probable’. But in the interests of ‘balancing’ our hash table, we’d like the hash codes $0, \dots, m - 1$ to be all equally likely.

Bad choice: $\#(s) = i(s) \bmod 128$. Effectively just last character of s .
So avoid powers of two!

Also not great: $\#(s) = i(s) \bmod 127$. Gives $\#(s) = \#(t)$ whenever s, t are **anagrams**. So $\#(\text{'algorithms'}) = \#(\text{'logarithms'})$.

Better: $\#(s) = i(s) \bmod 97$. Primes not too close to powers of two are reasonable.

Just the start of the delicate art of hash function design. . .

But whatever we do, **worst case** (all keys hashing to same code) is always terrible. A **malicious user** who knew your hash function could force this to happen . . .

Open addressing and probing

Solution 2: Rather than keeping bucket lists outside the hash table, store all items within the table itself (**open addressing**).

To deal with clashes, we use not just a simple hash function $\#(k)$, but a function $\#(k, i)$ where $0 \leq i < m$. For a key k :

- ▶ $\#(k, 0)$ is our **first choice** of hash value,
- ▶ $\#(k, 1)$ is our **second choice**, etc.

so that $\#(k, 0), \#(k, 1), \dots, \#(k, m - 1)$ is a **permutation** of $0, \dots, m - 1$. (Ideally, for a randomly chosen k , all $m!$ permutations should be equally likely.)

To **insert** an item e with key k , probe $A[\#(k, 0)], A[\#(k, 1)], \dots$ until we find a free slot $A[\#(k, i)]$, then put e there.

To **lookup** an item with key k , probe $A[\#(k, 0)], A[\#(k, 1)], \dots$ until we find either an item with key k , or free cell (lookup failed).

Probing: example

Let's use an array A of size $m = 10$ to store a set of integers.

0	1	2	3	4	5	6	7	8	9
58								28	49

Probe function: $\#(k, i) = (k + i) \bmod 10$.

insert(49). $\#(49, 0) = 9$: free.

insert(28). $\#(28, 0) = 8$: free.

insert(58). $\#(58, 0) = 8$: taken. $\#(58, 1) = 9$: taken.
 $\#(58, 2) = 0$: free.

contains(28). $\#(28, 0) = 8$, $A[8] = 28$. So true.

contains(58). $\#(58, 0) = 8$, $A[8] = 28 \neq 58$.

$\#(58, 1) = 9$, $A[9] = 49 \neq 58$.

$\#(58, 2) = 0$: $A[0] = 58$. So true.

contains(39). $\#(39, 0) = 9$, $A[9] = 49 \neq 39$.

$\#(39, 1) = 0$, $A[0] = 58 \neq 39$.

$\#(39, 2) = 1$, $A[1]$ free. So false.

Probing: pros and cons

- ▶ **Expected** number of probes for **insert** (and hence for **lookup**) stays low until table is nearly full. (Can show it's $1/(1 - \alpha)$ for unsuccessful lookup; less for successful one.)
- ▶ No need for pointers. The memory this saves can be 'spent' on increasing table size m and so decreasing load α ...
So compared to bucket lists, get faster **lookup** for same amount of memory.
- ▶ However, **delete** is a pain for the probing approach.
- ▶ Design of probing functions is again a delicate art (*linear probing, quadratic probing, double hashing, ...*).

See CLRS 11.4 for more details.

Radical alternative: Perfect hashing

- ▶ *All* the approaches we've mentioned are bad in the worst case: size of bucket/sequence of probes can be of length n .
- ▶ Even in typical cases, probably *some* buckets will be large relative to α . (Birthday paradox!)

If we could **avoid clashes altogether**, these problems would vanish!
Would get worst-case $\Theta(1)$ lookup.

If set of keys is **static** (no **insert/delete** required), may be worth finding a **perfect hash function** (no clashes) for this set of keys.

As part of **Coursework 1**, we'll explore a state-of-the-art approach to perfect hashing.

Reading:

CLRS Chapter 11. You can omit the theorems and their proofs, except for Theorem 11.1 which corresponds to slide 9.

Today's music:

Nikolai Rimsky-Korsakov, *Flight of the bumblebee*.
(Piano version by Sergei Rachmaninov.)