

Introduction to Algorithms and Data Structures

Lecture 15: DFS and graph structure

Mary Cryan

School of Informatics
University of Edinburgh

Recursive DFS (no explicit Stack)

Algorithm $\text{dfs}(G)$

1. Initialise Boolean array *visited*, setting all to FALSE
2. **for all** $v \in V$ **do**
3. **if** $\text{visited}[v] = \text{FALSE}$ **then**
4. $\text{dfsFromVertex}(G, v)$

Algorithm $\text{dfsFromVertex}(G, v)$

1. $\text{visited}[v] \leftarrow \text{TRUE}$
2. **for all** w adjacent to v **do**
3. **if** $\text{visited}[w] = \text{FALSE}$ **then**
4. $\text{dfsFromVertex}(G, w)$

- ▶ The recursive set-up eliminates the need for an explicit stack;
- ▶ We will have reversed prioritisation of the vertices adjacent to v .

Analysis of DFS

Lemma

During $\text{dfs}(G)$, $\text{dfsFromVertex}(G, v)$ is invoked exactly once for each vertex v .

Proof.

At least once:

- ▶ $\text{visited}[v]$ can only become `TRUE` when $\text{dfsFromVertex}(G, v)$ is executed.
- ▶ If $\text{visited}[v]$ remains `FALSE`, $\text{dfsFromVertex}(G, v)$ will eventually be called by line 4 of $\text{dfs}(G)$.

At most once:

- ▶ First call of $\text{dfsFromVertex}(G, v)$ sets $\text{visited}[v]$ to `TRUE`.
- ▶ After $\text{visited}[v]$ is `TRUE`, $\text{dfsFromVertex}(G, v)$ is *never* called again.



(“At most once” is also true for Stack dfs, but “at least once” is not. dfsFromVertex is more to “start a component” in the Stack version)

Analysis of DFS (cont'd)

Lemma

For a directed graph, $\sum_{v \in V} \text{out-degree}(v) = m$.

For an undirected graph, $\sum_{v \in V} \text{deg}(v) = 2m$.

Proof.

Every edge is counted exactly once on both sides of the equation (for directed).

For the undirected case, every edge is counted twice on the lhs. □

Analysis of recursive DFS

$G = (V, E)$ graph with n vertices and m edges

Algorithm dfs(G)

1. Initialise Boolean array *visited*, setting all to FALSE
2. **for all** $v \in V$ **do**
3. **if** *visited*[v] = FALSE **then**
4. dfsFromVertex(G, v)

- ▶ dfs(G): Ignoring calls to dfsFromVertex, total time $\Theta(n)$
- ▶ dfsFromVertex(v) is called at most once for every vertex v , and does $\Theta(\text{out-degree}(v))$ work, excluding recursive calls.

Overall time:

$$\begin{aligned}T(n, m) &= \Theta(n) + \sum_{v \in V} \Theta(\text{out-degree}(v)) \\&= \Theta\left(n + \sum_{v \in V} \text{out-degree}(v)\right) \\&= \Theta(n + m)\end{aligned}$$

Adjacency List or Adjacency Matrix?

We said each call to `dfsFromVertex(v)` takes $\Theta(\text{out-degree}(v))$ time (excluding recursive calls).

Algorithm `dfsFromVertex(G, v)`

1. *visited*[v] \leftarrow TRUE
2. **for all** w adjacent to v **do**
3. **if** *visited*[w] = FALSE **then**
4. `dfsFromVertex(G, w)`

If we are iterating over “all w adjacent to v ” in $\Theta(\text{out-degree}(v))$ time, then we **must** be using an Adjacency list structure.

Analysis of original DFS

Compare the two `dfsFromVertex(G, v)` methods:

Algorithm `dfsFromVertex(G, v)`

1. $visited[v] \leftarrow \text{TRUE}$
2. **for all** w adjacent to v **do**
3. **if** $visited[w] = \text{FALSE}$ **then**
4. `dfsFromVertex(G, w)`

Algorithm `dfsFromVertex(G, v)`

1. `S.push(v)`
2. **while not** `S.isEmpty()` **do**
3. $u \leftarrow S.pop()$
4. **if** $visited[u] = \text{FALSE}$ **then**
5. $visited[u] \leftarrow \text{TRUE}$
6. **for all** w adjacent to u
7. `S.push(w)`

$visited[v] \leftarrow \text{TRUE}$ <code>dfsFromVertex(v)</code>	\leftrightarrow	<code>S.pop();</code> $visited[v] \leftarrow \text{TRUE}$ for all w adjacent to v ; <code>S.push(w); while</code>
--	-------------------	---

Recursive: marks v as “visited”, *then* calls `dfsFromVertex` for unvisited adjacent vertices

Iterative: marks v as “visited” *after* “popping” it, then “pushes” all adjacent vertices. However, the number of Stack operations for v is bounded in terms of *the number of edges into* $v \Rightarrow$ the overall runtime for our original dfs is still $\Theta(n + m)$.



DFS Forests

A DFS traversing a graph builds up a *forest* whose vertices are all vertices of the graph and whose edges are all vertices traversed during the DFS.

Definition

A vertex w is a *child* of a vertex v in the DFS forest if $\text{dfsFromVertex}(G, w)$ is called from $\text{dfsFromVertex}(G, v)$.

DFS Forests Example



Recall Q2 of tutorial sheet 5 on how the connected components can vary depending on the order in which we consider vertices at the top-level of dfs.

Topological Sorting

Example:

10 tasks to be carried out. Some of them depend on others.

- ▶ Task 0 must be completed before Task 1 can be started.
- ▶ Task 1 and Task 2 must be done before Task 3 can start.
- ▶ Task 4 must be done before Task 0 or Task 2 can start.
- ▶ Task 5 must be done before Task 0 or Task 4 can start.
- ▶ Task 6 must be done before Task 4, 5 or 7 can start.
- ▶ Task 7 must be done before Task 0 or Task 9 can start.
- ▶ Task 8 must be done before Task 7 or Task 9 can start.
- ▶ Task 9 must be done before Task 2 or Task 3 can start.


Topological order

Definition

Let $G = (V, E)$ be a directed graph. A *topological order* of G is a total order \prec of the vertex set V such that for all edges $(v, w) \in E$ we have $v \prec w$.

(in some fields this is called a *linear extension*)

Tasks as a (directed) graph



Does this graph have a topological order?

Yes. One topological sort is:

$8 \prec 6 \prec 7 \prec 9 \prec 5 \prec 4 \prec 2 \prec 0 \prec 1 \prec 3.$

Topological order (cont'd)

A digraph that has a cycle does not have a topological order.

Definition

A *DAG* (**d**irected **a**cyclic **g**raph) is a digraph without cycles.

Theorem

A digraph has a topological order if and only if it is a DAG.

Classification of vertices during recursive DFS

$G = (V, E)$ graph, $v \in V$. Consider $\text{dfs}(G)$.

- ▶ v is **finished** if $\text{dfsFromVertex}(G, v)$ has been completed.

Vertices can be in the following states:

- ▶ not yet visited (let us call a vertex in this state *white*),
- ▶ visited, but not yet finished (*grey*).
- ▶ finished (*black*).

Classification of vertices during recursive DFS (cont'd)

Lemma

Let G be a graph and v a vertex of G . Consider the moment during the execution of $\text{dfs}(G)$ when $\text{dfsFromVertex}(G, v)$ is started.

Then for all vertices w we have:

- 1. If w is white and reachable from v , then w will be black before v .*
- 2. If w is grey, then v is reachable from w .*

Topological sorting

$G = (V, E)$ digraph. Define order on V as follows:

$$v \prec w \iff w \text{ becomes black before } v.$$

Theorem

If G is a DAG then \prec is a topological order.

Proof.

Suppose $(v, w) \in E$. Consider $\text{dfsFromVertex}(G, v)$.

- ▶ If w is already *black*, then $v \prec w$ (and this is what we want).
- ▶ If w is *white*, then by Lemma part 1., w will be *black* before v . Thus $v \prec w$.
- ▶ If w is *grey*, then by Lemma part 2. v is reachable from w . So there is a path p from w to v . Path p and edge (v, w) together form a cycle.

Contradiction! (G is acyclic ...)



Topological sorting implemented

Algorithm topSort(G)

1. Initialise array *state*
by setting all entries to *white*.
2. Initialise linked list L
3. **for all** $v \in V$ **do**
4. **if** $state[v] = white$ **then**
5. sortFromVertex(G, v)
6. **print** L

Topological sorting implemented


Algorithm sortFromVertex(G, v)

1. $state[v] \leftarrow grey$
2. **for all** w adjacent to v **do**
3. **if** $state[w] = white$ **then**
4. sortFromVertex(G, w)
5. **else if** $state[w] = grey$ **then**
6. **print** “ G has a cycle”
7. **halt**
8. $state[v] \leftarrow black$
9. $L.insertFirst(v)$

Difference from dfs itself - the order the vertices get added to the list.

Running-time is again $\Theta(n + m)$

Example



Use the algorithm topSort to compute a topological sort of this graph (video).

Connected components of an undirected graph

$G = (V, E)$ undirected graph

Definition

- ▶ A subset C of V is *connected* if for all $v, w \in C$ there is a path from v to w (if G is directed, say *strongly connected*).
- ▶ A *connected component* of G is a **maximum connected subset** C of V . (no connected subset C' of V strictly contains C).
- ▶ G is *connected* if it only has one connected component, that is, if for all vertices v, w there is a path from v to w .

Connected components - undirected (cont'd)

- ▶ Each vertex of an undirected graph is contained in exactly one connected component.
- ▶ For each vertex v of an undirected graph, the connected component that contains v is precisely the set of all vertices that are reachable from v .

For an **undirected** graph G , `dfsFromVertex(G, v)` visits exactly the vertices in the connected component of v .

And the same is true for `bfsFromVertex(G, v)` (either will do!)

Connected components - undirected (cont'd)

Algorithm connComp(G)

1. Initialise Boolean array *visited*
by setting all entries to FALSE
2. **for all** $v \in V$ **do**
3. **if** $visited[v] = \text{FALSE}$ **then**
4. **print** "New Component"
5. ccFromVertex(G, v)

Algorithm ccFromVertex(G, v)

1. $visited[v] \leftarrow \text{TRUE}$
2. **print** v
3. **for all** w adjacent to v **do**
4. **if** $visited[w] = \text{FALSE}$ **then**
5. ccFromVertex(G, w)

Reading

From [CLRS] as usual:

- ▶ Computing topological sort - Section 22.4

Hope you get a break over the holidays!

And “see” you in 2022.