


Intelligent Agents and their Environments

Petros Papapanagiotou

Informatics 2D: Reasoning and Agents

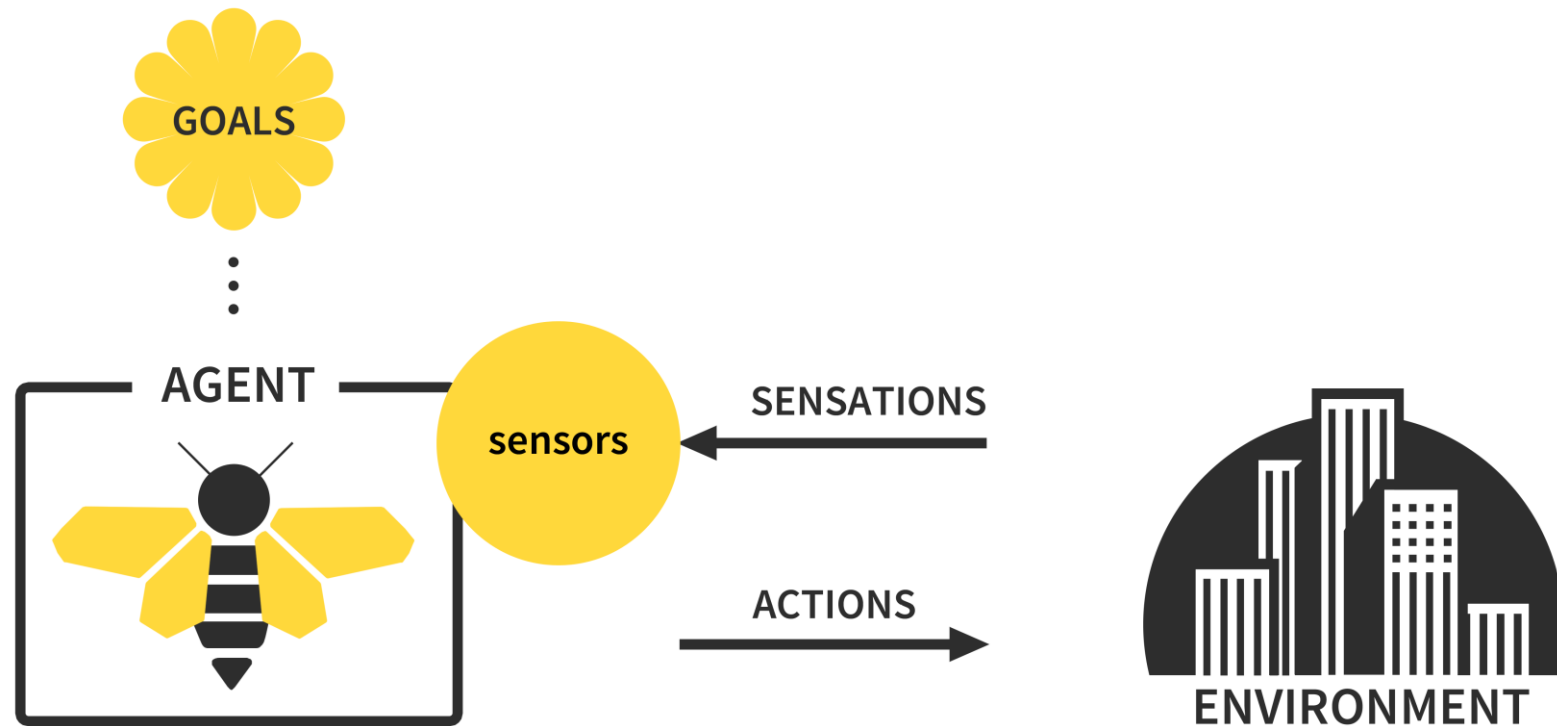
Lecture 1



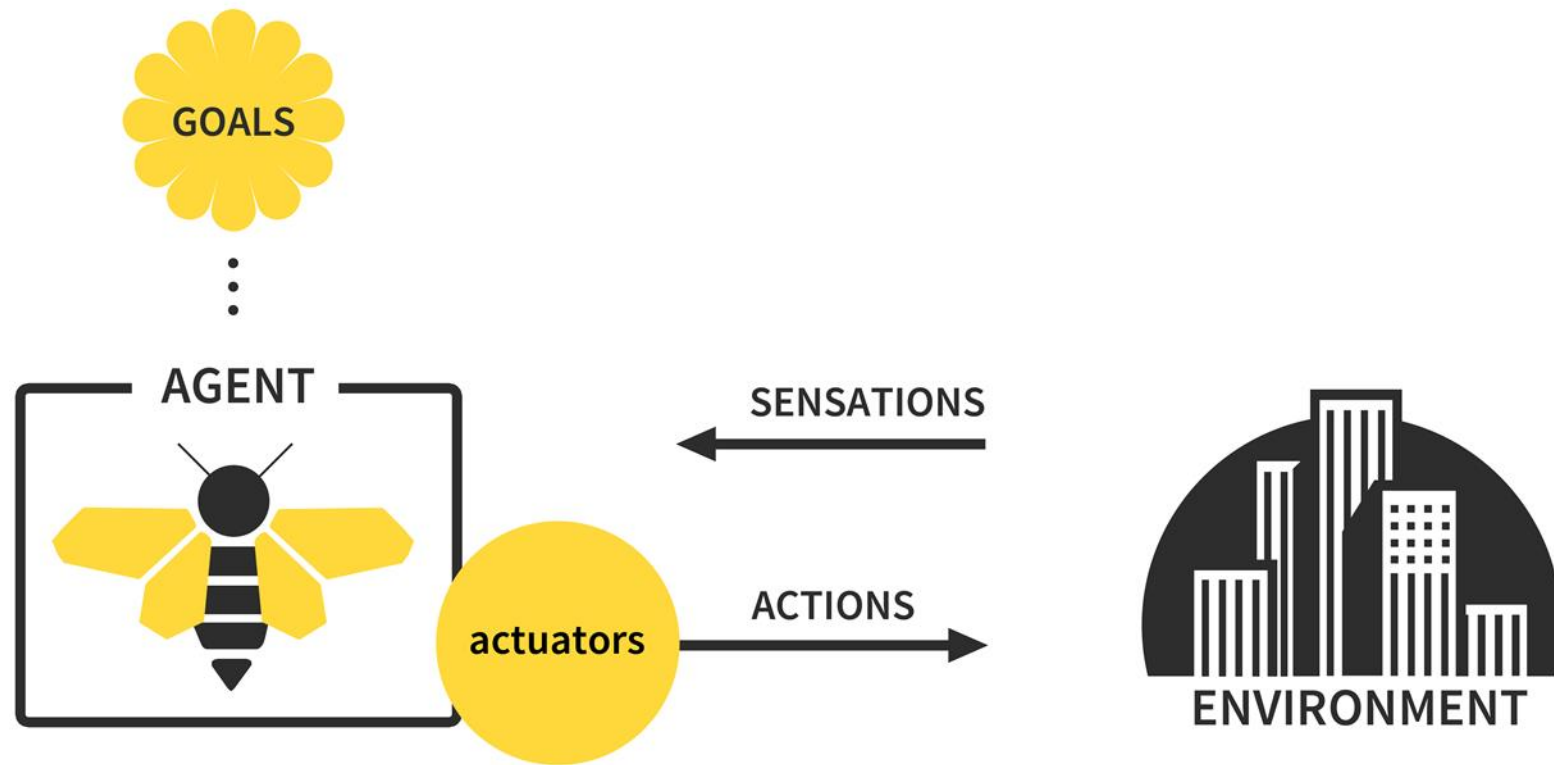
What is an Intelligent Agent?



Structure



Structure



An agent

Perceives its *environment*,
Through its *sensors*,
Then achieves its *goals*,
By acting on its environment via *actuators*.

Categorise agents



Environment



Goals



Percepts



Actions

Example: Mail sorting



Conveyor belt of letters



Route letter into correct bin



Array of pixel intensities



Route letter into bin





18.1 °C
Living Room
Temperat...

20.5 °C
Main Bedroom
Temperat...

20.2 °C
Bedroom
Temperat...

20.6 °C
Kitchen
Temperat...

20.5 °C
Bathroom
Temperat...

golde...
Daylight


Sun


75.9
kane-NAS
Volume
Used


Main Door






Main Door
Closed

Battery



1+7Pro

INF2D: REASONING AND AGENTS



Christmas Lights



Living Room Lights



Kitchen Lights



Baby Lights

Smart Home

Example: Smart home



occupants enter and leave house, occupants enter and leave rooms; daily variation in outside light and temperature



occupants warm, room lights are on when room is occupied, house energy efficient



signals from temperature sensors, movement sensors, clock, sound sensor, weather sensor



room heaters on/off, lights on/off



Example: Autonomous car



streets, other vehicles, pedestrians, traffic signals/lights/signs



safe, fast, legal trip



camera, GPS signals, speedometer, sonar



steer, accelerate, brake



Type of Intelligent Agents

Simple Reflex Agents

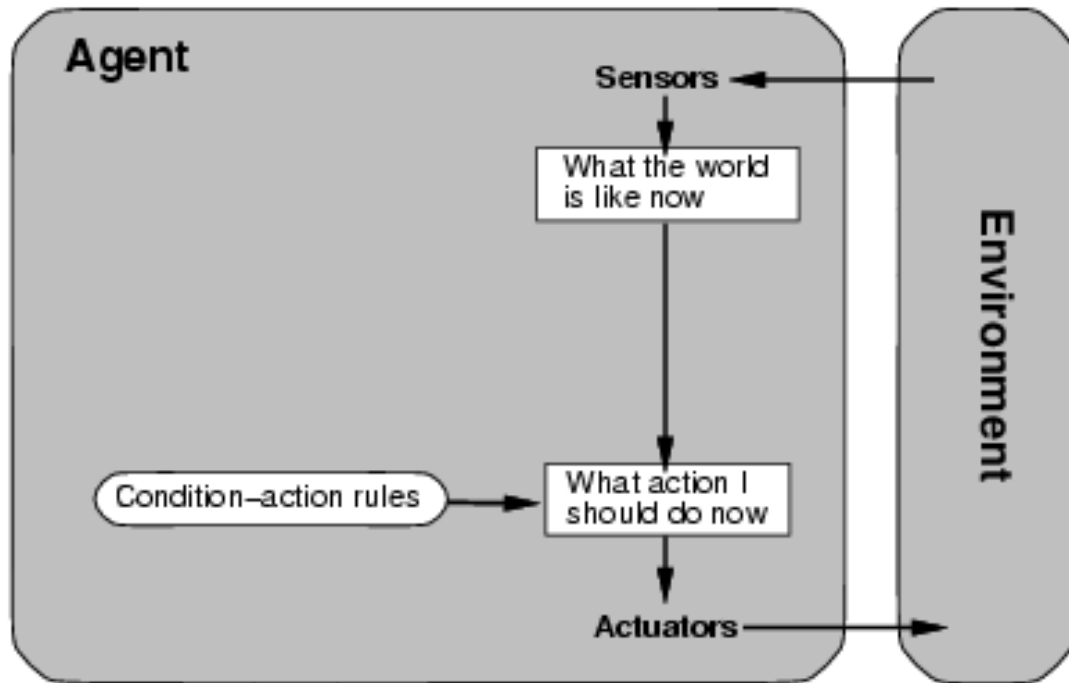
Action depends only on immediate percepts.

Implement by *condition-action rules*.

Example:

- *Agent*: Mail sorting robot
- *Environment*: Conveyor belt of letters
- *Rule*: e.g. *city=Edinburgh* → *put in Scotland bag*

Simple Reflex Agents



```
function SIMPLE-REFLEX-AGENT(percept)
  returns action
  persistent: rules (set of condition-action rules)
    state  $\leftarrow$  INTERPRET-INPUT(percept)
    rule  $\leftarrow$  RULE-MATCH(state, rules)
    action  $\leftarrow$  rule.ACTION
return action
```

Model-Based Reflex Agents

Action may depend on history or unperceived aspects of the world.

Need to maintain *internal world model*.

Example:

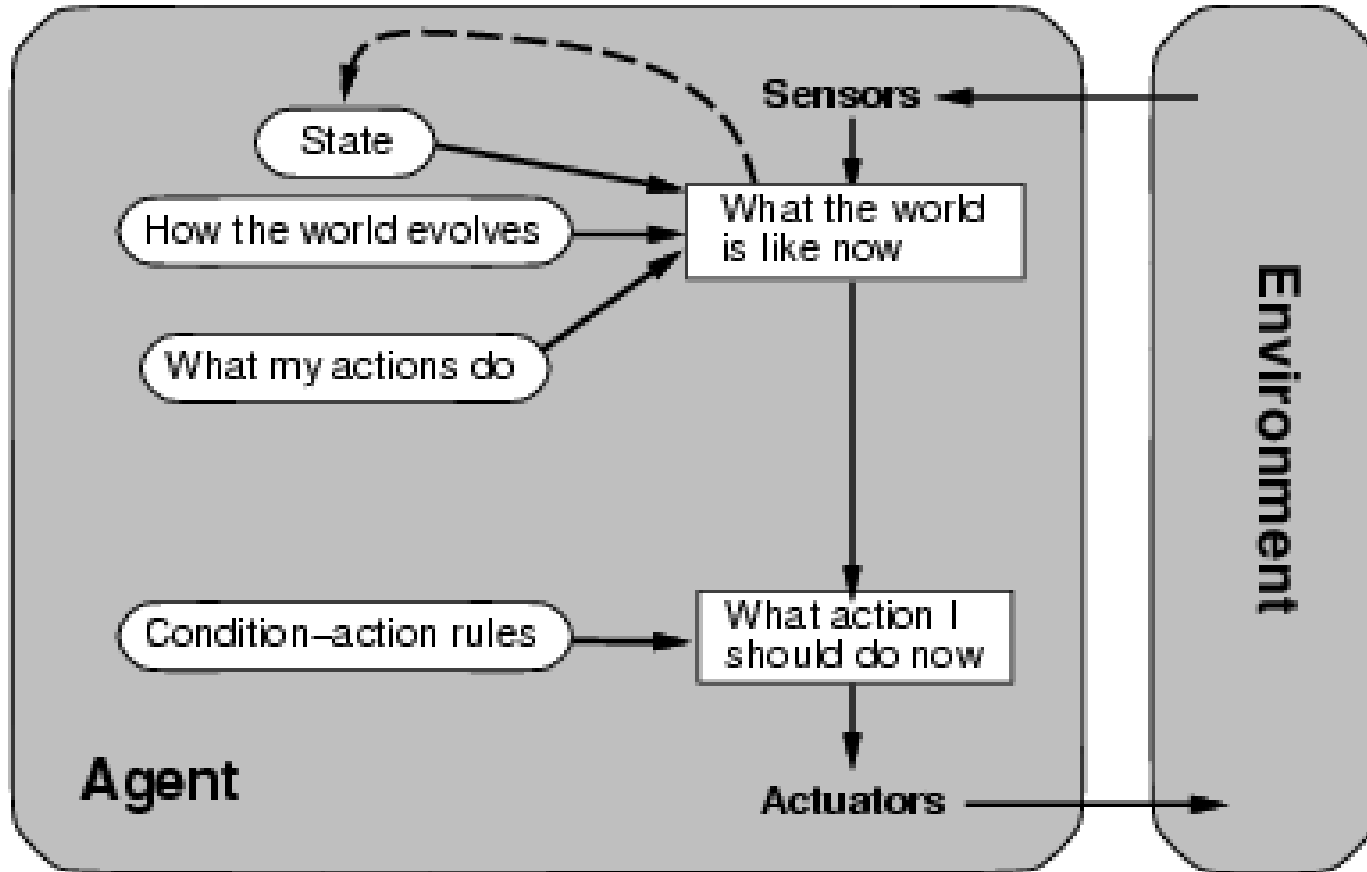
Agent: robot vacuum cleaner

Environment: dirty room, furniture

Model: map of room, which areas already cleaned

Sensor/model tradeoff.

Model-Based Reflex Agents



```
function REFLEX-AGENT-WITH-STATE(percept)  
returns action  
persistent: state, description of current world state  
              model, description of how the next state  
                  depends on current state and action  
              rules, a set of condition-action rules  
              action, the most recent action, initially none  
state ← UPDATE-STATE(state, action, percept, model)  
rule ← RULE-MATCH(state, rules)  
action ← rule.ACTION  
return action
```

Goal-Based Agents

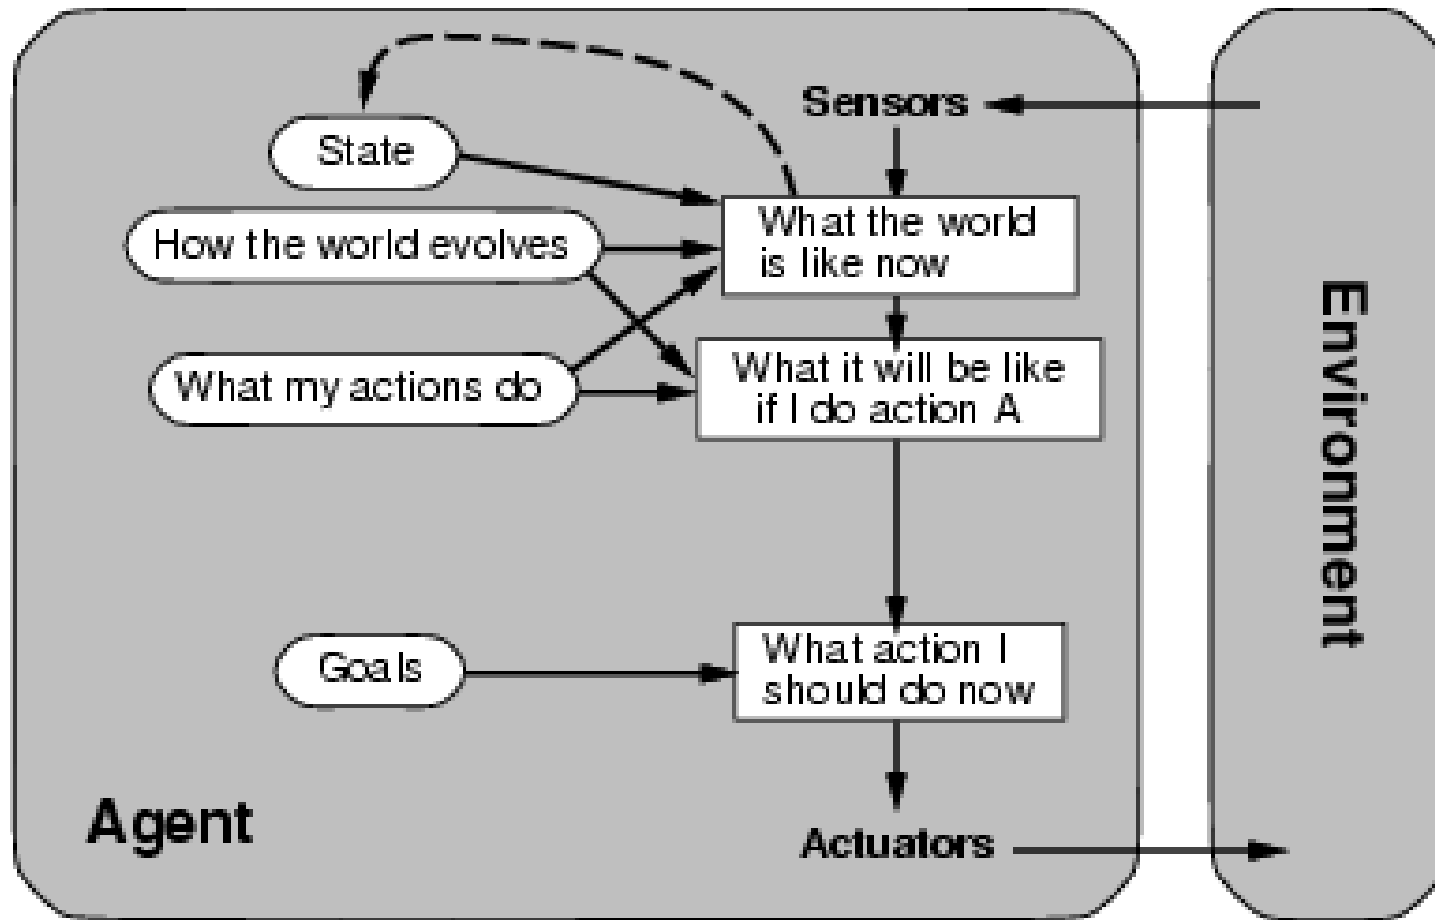
Agents so far have fixed, implicit goals.

We want agents with variable goals.

Forming plans to achieve goals is a topic for later.

Example:

- **Agent**: robot maid
- **Environment**: house & people.
- **Goals**: clean clothes, tidy room, table laid, etc



Goal-Based Agents

Utility-Based Agents

Agents so far have had a single goal.

Agents may have to juggle conflicting goals.

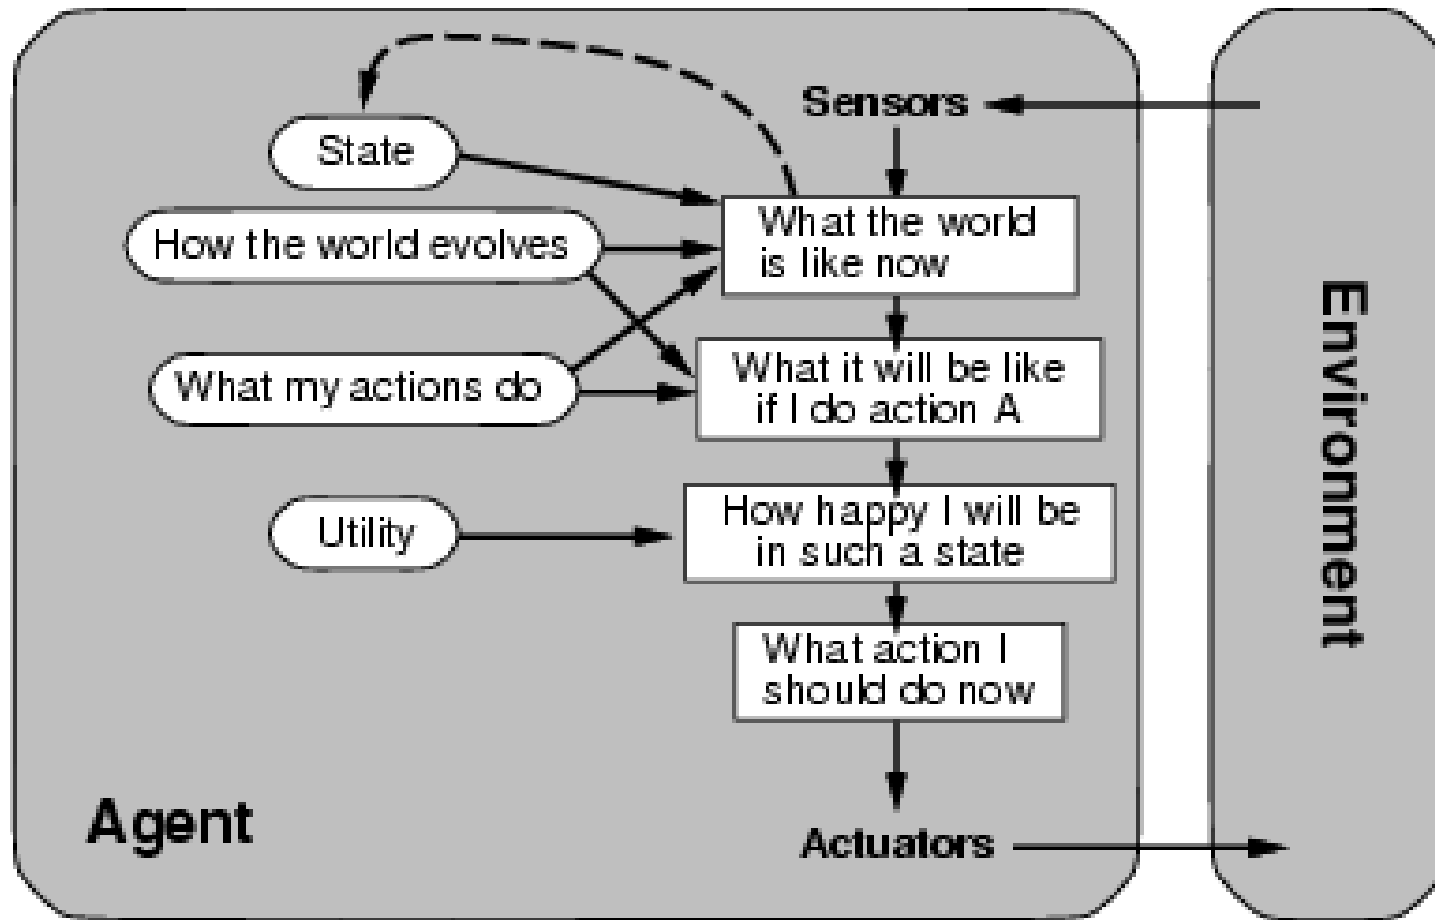
Need to optimise utility over a range of goals.

Utility: measure of *goodness* (a real number).

Combine with probability of success to get *expected utility*.

Example:

- **Agent**: autonomous car.
- **Environment**: roads, vehicles, signs, etc.
- **Goals**: stay safe, reach destination, be quick, obey law, save fuel, etc.



Utility-Based Agents

Learning Agents

How do agents improve their performance in the light of experience?

- Generate problems which will test performance.
- Perform activities according to rules, goals, model, utilities, etc.
- Monitor performance and identify non-optimal activity.
- Identify and implement improvements

We will not be covering learning agents, but this topic is dealt with in several honours-level courses (see also R&N, Ch. 18-21).

Exercise

Consider a chess playing program.

What sort of agent would it need to be?



Solution

Simple-reflex agent: but some actions require some memory (e.g. [castling](#))

Model-based reflex agent: but needs to reason about future

Goal-based agent: but only has one goal

Utility-based agent: might consider multiple goals with limited lookahead

Learning agent: learns from experience or self-play



Types of Environments

Observable?

FULLY



[Source](#)

PARTIALLY



[Source](#)

Observable?

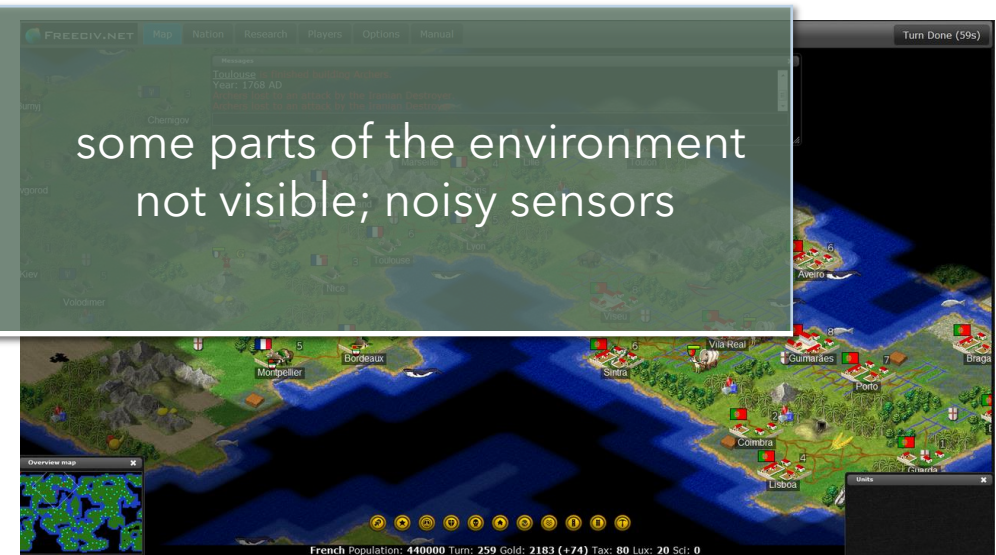
FULLY



[Source](#)

vs.

PARTIALLY



[Source](#)

Deterministic?

DETERMINISTIC




STOCHASTIC



[Source](#)

Deterministic?

DETERMINISTIC



next state fully determined by
current state and agent's actions

vs.

STOCHASTIC



random changes - cannot be
predicted exactly

An environment may appear stochastic if it is only partially observable.

[Source](#)

Sequential?

EPISODIC



[Source](#)

SEQUENTIAL



Sequential?


EPISODIC



next episode does not depend
on previous actions

vs.

SEQUENTIAL




actions affect the future

[Source](#)

Static?

STATIC



DYNAMIC




[Source](#)

Static?

STATIC

environment unchanged while
agent deliberates



vs.

DYNAMIC

environment can change at any
time or even continuously



[Source](#)

Discrete?

DISCRETE



[Source](#)

CONTINUOUS



[Source](#)

Discrete?


DISCRETE



[Source](#)

vs.


CONTINUOUS



[Source](#)

How many agents?

SINGLE




MULTI-AGENT



[Source](#)

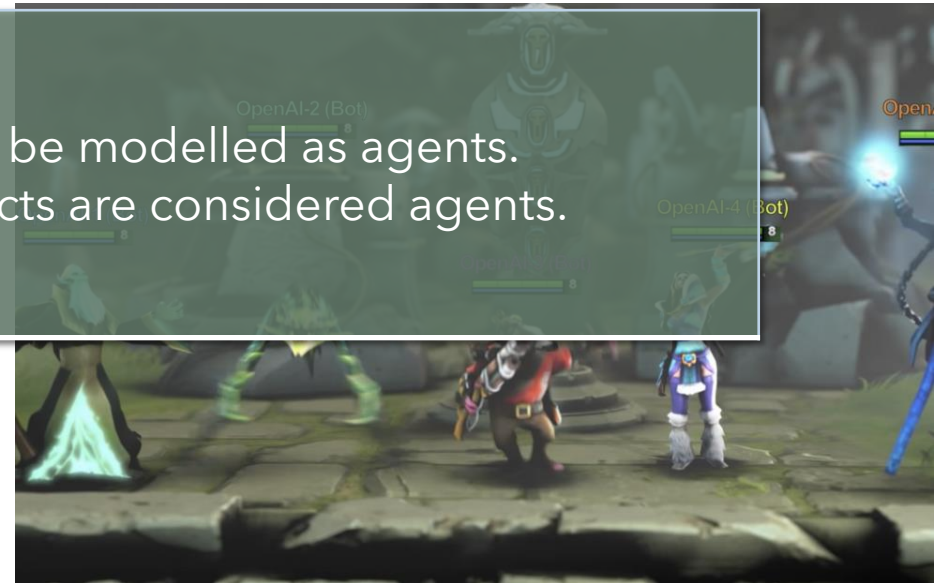
How many agents?

SINGLE



Consider how many object must be modelled as agents.
Element of choice over which objects are considered agents.

MULTI-AGENT



[Source](#)

Summary

An agent might have any combination of these properties:

- from “**benign**”: i.e., fully observable, deterministic, episodic, static, discrete and single agent
- to “**chaotic**”: partially observable, stochastic, sequential, dynamic, continuous and multi-agent

What are the properties of the environment that would be experienced by

- a mail-sorting robot?
- a smart home?
- a car-driving robot?

Why?

- Understanding what is an agent and how it can be modelled.
- Understanding the environment and the assumptions or considerations that need to be made.
- Making the right design decisions and choosing the right tools.
- Managing the complexity.