


# Adversarial Search

---

*Petros Papapanagiotou*

Informatics 2D: Reasoning and Agents

**Lecture 8**



# Games vs. search problems

---

"Unpredictable" opponent → solution is a **strategy** / **policy**

- Specify a move for *every possible* opponent reply

Time limits → unlikely to find goal, must approximate



TYPES OF GAMES	deterministic	chance
perfect information	chess, checkers	backgammon, monopoly
imperfect information	battleships, stratego	bridge, poker, scrabble

# Games vs. search problems


---

We are interested in zero-sum games of perfect information:

- Deterministic, fully observable
- Agents act alternately
- Utilities at end of game are equal and opposite (adding up to 0)

# Game tree (2-player, deterministic, turns)

- 2 players: MAX and MIN
- MAX moves first
- Tree built from MAX's POV



# Optimal Decisions

---

## Normal search:

- optimal decision is a *sequence of actions* leading to a goal state (i.e. a winning terminal state)

## Adversarial search:

- MIN has a say in game
- MAX needs to find a contingent strategy which specifies:
  - MAX's **move** in initial state then...
  - MAX's **moves** in states resulting from every response by MIN to the **move** then...
  - MAX's moves in states resulting from every response by MIN to all those **moves**, etc...

# Minimax value

---

minimax value of a node = utility for MAX of being in corresponding state:

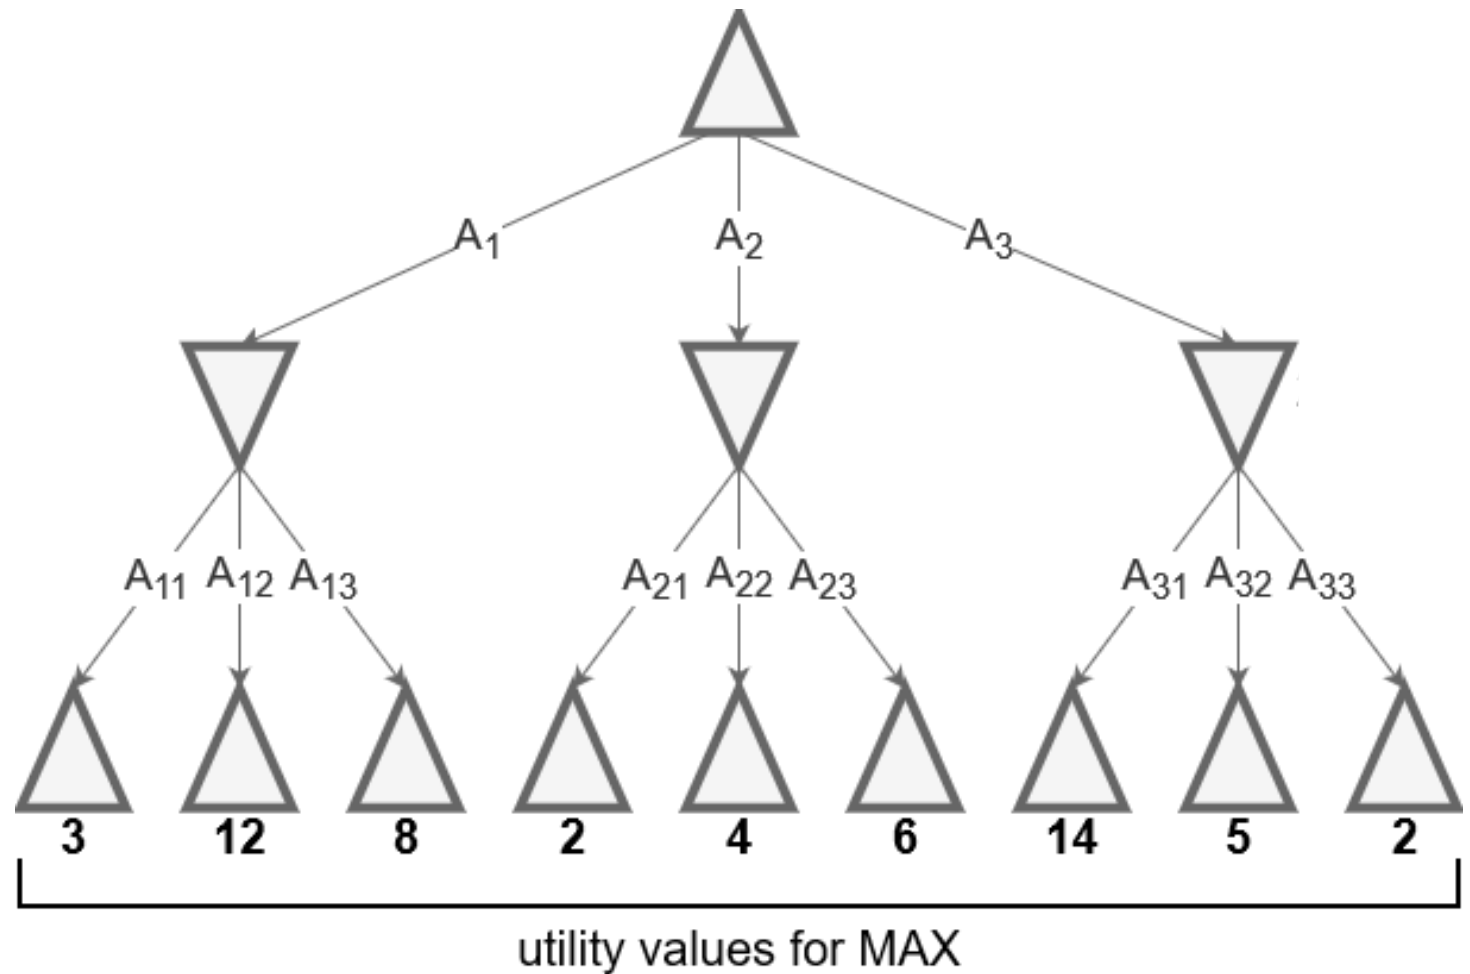
$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s,a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s,a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

# Minimax

Perfect play for  
deterministic, perfect-  
information games

Idea: choose move to  
position with highest  
minimax value

= best achievable payoff  
against best play

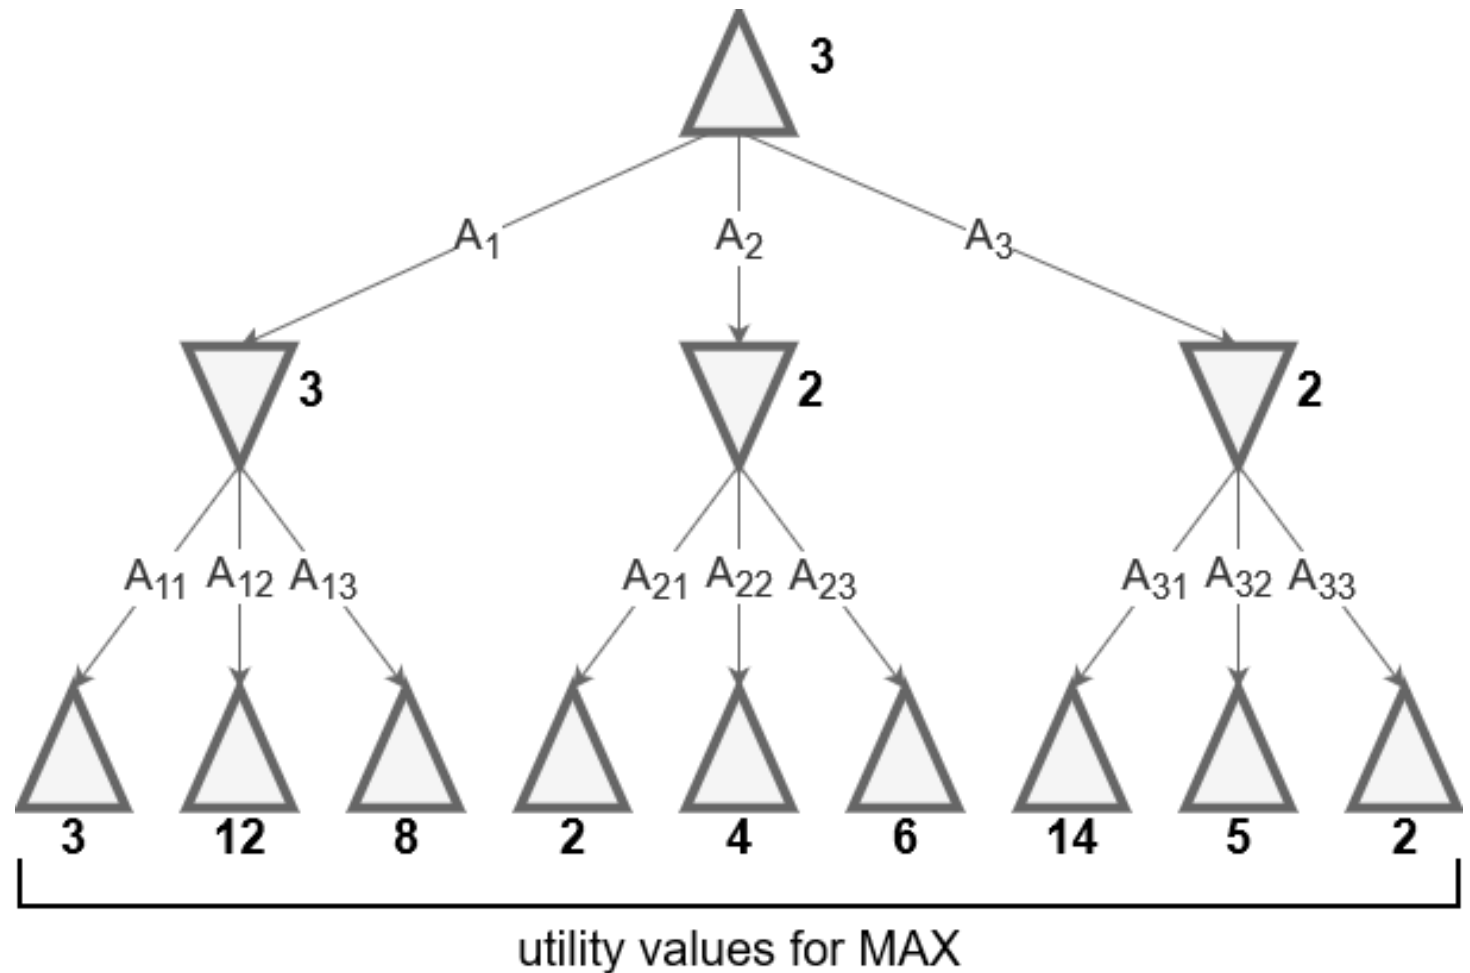


# Minimax

Perfect play for  
deterministic, perfect-  
information games

Idea: choose move to  
position with highest  
minimax value

= best achievable payoff  
against best play





# Minimax algorithm

---

Idea:

- Proceed all the way down to the leaves of the tree
- then minimax values are backed up through tree

```
function MINIMAX-DECISION(state) returns an action  
  return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$ 
```

---

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$   
  return v
```

---

```
function MIN-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$   
  return v
```

# Properties of Minimax

---




**Complete?**



**Time complexity?**



**Space complexity?**



**Optimal?**

# Properties of Minimax

---



**Complete?**


Yes (if tree is finite)



**Time complexity?**



**Space complexity?**



**Optimal?**

# Properties of Minimax

---



**Complete?**

Yes (if tree is finite)




**Time complexity?**

$O(b^m)$



**Space complexity?**



**Optimal?**

# Properties of Minimax

---



## **Complete?**

Yes (if tree is finite)




## **Time complexity?**

$O(b^m)$



## **Space complexity?**

$O(bm)$



## **Optimal?**

# Properties of Minimax

---



## **Complete?**

Yes (if tree is finite)




## **Time complexity?**

$O(b^m)$



## **Space complexity?**

$O(bm)$



## **Optimal?**

Yes (against an optimal opponent)

# Complexity

---

For chess,  $b \approx 35$ ,  $m \approx 100$  (average  $\approx 40$ ) for "reasonable" games

- exact solution completely infeasible!
- would like to eliminate (large) parts of game tree

$$35^{40} = 5.791 \cdot 10^{61}$$

$$35^{100} = 2.552 \cdot 10^{154}$$


# $\alpha$ - $\beta$ pruning

---




# $\alpha$ - $\beta$ pruning example

---




# $\alpha$ - $\beta$ pruning example

---




# $\alpha$ - $\beta$ pruning example

---




# $\alpha$ - $\beta$ pruning example

---




# $\alpha$ - $\beta$ pruning example

---



# $\alpha$ - $\beta$ pruning example

---




# $\alpha$ - $\beta$ pruning example

Are minimax value of root and, hence, minimax decision *independent* of pruned leaves?

Let pruned leaves have values  $u$  and  $v$ ,

$$\begin{aligned}\text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, u, v), \min(14, 5, 2)) \\ &= \max(3, \min(2, u, v), 2) \\ &= \max(3, z, 2) \quad \text{where } z \leq 2 \\ &= 3\end{aligned}$$

Yes!



MAX

MIN


..

..

..

MAX

MIN



# Why is it called $\alpha$ - $\beta$ ?

$\alpha$  is the value of the best (i.e., highest-value) choice found so far at any choice point along the path for MAX

If  $v$  is worse than  $\alpha$ , MAX will avoid it  
→ prune that branch

$\beta$  is defined symmetrically for MIN



# The $\alpha$ - $\beta$ algorithm

---

$\alpha$  is value of the best i.e. **highest**-value choice found so far at any choice point along the path for MAX

$\beta$  is value of the best i.e. **lowest**-value choice found so far at any choice point along the path for MIN

```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value  $v$ 
```

---

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

---

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```

# Complexity of $\alpha$ - $\beta$

---

Pruning does not affect final result (as we saw for example)

Good move ordering improves effectiveness of pruning

With “perfect ordering”, time complexity =  $O(b^{m/2})$

- branching factor goes from  $b$  to  $\sqrt{b}$
- **doubles solvable depth** of search compared to minimax

A simple example of the value of reasoning about which computations are relevant (a form of meta-reasoning)

# Resource limits

---

# Resource limits

---

Suppose we have 100 secs and can explore  $10^4$  nodes/sec

➤  $10^6$  nodes per move

➤  $b^m = 10^6$

➤ For  $b = 35 \rightarrow 35^4 = 1.5 \cdot 10^6 \rightarrow$  so  $m \approx 4$

4-ply lookahead is a hopeless chess player!

- 4-ply  $\approx$  human novice
- 8-ply  $\approx$  typical PC, human master
- 12-ply  $\approx$  Deep Blue, Kasparov

# Standard approach

---

## Cutoff test

e.g., depth limit (perhaps add **quiescence search**, which tries to search interesting positions to a greater depth than quiet ones)

## Evaluation function

= estimated desirability of position

# Standard approach

*MinimaxCutoff* is identical to *MinimaxValue* except:

1. *TERMINAL-TEST* is replaced by *CUTOFF*
2. *UTILITY* is replaced by *EVAL*

```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

---

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

---

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```

# Evaluation functions

---

Often a linear weighted sum of **features**

$$\text{EVAL}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

where each  $w_i$  is a weight and each  $f_i$  is a feature of state  $s$

## Chess example

- queen = 1, king = 2, etc.
- $f_i$  = number of pieces of type  $i$  on board
- $w_i$  = value of the piece of type  $i$

# Deterministic games in practice

---





# Checkers

---

Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a precomputed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.



<http://jonathanschaeffer.blogspot.com/2012/08/chinook-twenty-years-later.html>

# Othello

---

Human champions refuse to compete against computers.

Logistello, written by Michael Buro, defeated the human world champion Takeshi Murakami six games to none in 1997.

The best Othello programs are now much stronger than any human player.

	a	b	c	d	e	f	g	h	
1	●	●	●	●	●	●	●	●	1
2	●	●	○	●	●	●	●	●	2
3	●	●	●	○	●	●	○	●	3
4	●	●	○	●	●	○	●	●	4
5	●	●	●	●	●	●	●	●	5
6	●	●	●	○	●	●	●	●	6
7	●	●	○	○	○	●	●	●	7
8	●	●	●	●	●	●	●	●	8
	a	b	c	d	e	f	g	h	

<https://skatgame.net/mburo/log.html>



# Chess

---

Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40-ply.

[https://en.wikipedia.org/wiki/Deep\\_Blue\\_\(chess\\_computer\)](https://en.wikipedia.org/wiki/Deep_Blue_(chess_computer))

# Modern Chess



## Stockfish

- Uses an advanced version of  $\alpha$ - $\beta$  pruning among other algorithms.
- Recently added a simple neural network in its evaluation.
  - Improved by 100+ ELO points since.
- Analyses  $10^8$  positions per second (half when using the neural network).

## AlphaZero (successor of AlphaGo Zero)

- Based on Monte Carlo tree search, deep neural networks and self-play.
- Analyses 80,000 positions per second.
- Defeated Stockfish with 28W-72D-0L in 2016.

## Leela Zero

- Released 2017 with ideas from AlphaGo Zero's paper.
- Believed to have surpassed AlphaZero.
- Neck to neck with modern Stockfish, losing narrowly to it in the last 3 TCEC superfinals.

# Go

---


Human champions used to refuse to compete against computers, because they were are too bad.

In Go,  $b > 300$ , so most programs use pattern knowledge bases to suggest plausible moves.

In 2015 AlphaGo became the first computer program to beat a human professional Go player (Fan Hui) without handicap.

In 2016 AlphaGo beat world's #2 Lee Sedol 4-1.

Evolved into AlphaGo Zero (without human datasets), then AlphaZero, and more recently MuZero (model-free) .



Game 4, Lee Sedol (white) v. AlphaGo (black).  
First 78 moves

[https://en.wikipedia.org/wiki/Lee\\_Sedol](https://en.wikipedia.org/wiki/Lee_Sedol)

# Summary

---

Games are fun to work on!

They illustrate several important points about AI.

Perfection is unattainable → must approximate!

Good idea to think about what to think about.

Modern AI demonstrating superhuman performance.